MINSOO KIM

222, Wangsimni-ro, Seongdong-gu, Seoul, Republic of Korea, 04763

+82-10-8203-6871 marsjacobs.github.io minsoo2333@hanyang.ac.kr

RESEARCH INTERESTS

Efficient deep learning inference algorithm, LLM Quantization, Knowledge Distillation, LLM Fine-Tuning (PEFT)

EDUCATION

Ph.D. Candidate in Electronic Engineering Mar. 2021 - Present
Hanyang University, Seoul, South Korea
Advisor: Professor Jungwook Choi

B.S in Electronic Engineering Feb. 2021
Hanyang University, Seoul, South Korea

PUBLICATIONS

- [**NeurIPS 2023**] **Minsoo Kim**, Sihwa Lee, Jangwhan Lee, Sukjin Hong, Du-Seong Chang, Wonyong Sung and Jungwook Choi "Token-Scaled Logit Distillation for Ternary Weight Generative Language Models", *Thirty-seventh Conference on Neural Information Processing Systems*
[Paper]
- [**EMNLP 2023 main**] Janghwan Lee*, **Minsoo Kim***, Seungcheol Baek, Seok Joong Hwang, Wonyong Sung and Jungwook Choi "Enhancing Computation Efficiency in Large Language Models through Weight and Activation Quantization", *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (to appear)
- [**EACL 2023 main**] **Minsoo Kim**, Kyuhong Shim, Seongmin Park, Wonyong Sung and Jungwook Choi, "Teacher Intervention: Improving Convergence of Quantization Aware Training for Ultra-Low Precision Transformers", *The 17th Conference of the European Chapter of the Association for Computational Linguistics*
[Paper, Code Poster]
- [**EMNLP 2022 main**] **Minsoo Kim**, Sihwa Lee, Sukjin Hong, Du-Seong Chang, and Jungwook Choi, "Understanding and Improving Knowledge Distillation for Quantization-Aware Training of Large Transformer Encoders," *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*
[Paper, Code, Poster]
- [**DAC 2022**] Joonsang Yu, Junki Park, Seongmin Park, **Minsoo Kim**, Sihwa Lee, Donghyun Lee, Jungwook Choi, "NN-LUT: neural approximation of non-linear operations for efficient transformer inference, " *Proceedings of the 59th ACM/IEEE Design Automation Conference*
[Paper]
- Hyeonseung Kim, **Minsoo Kim**, Jungwook Choi, "Improving training method for very low bit weight quantization of Light Deep Learning Model, " Autumn Annual Conference of IEIE 2020

RESEARCH EXPERIENCE

Research Assistant, AI Algorithm & Hardware Lab, Hanyang University — Mar 2021 - Present
Advisor: Professor Jungwook Choi — *Seoul, South Korea*

- **GPT based generative LLM compression & auto-regressive text generation operation analysis**
  - Analyze the biased word generation behavior in GPT-2 models under 2-bit weight quantization with knowledge distillation.
  - Propose new scaled KD method achieving comparable perplexity to Full-Precision teacher model with 2-bit weight quantized GPT-2/OPT model.
  - Profile the text generation inference workload in single GPU for GPT-2/3 models, identifying memory-bound and low-density computation challenges in GPU architecture with text-generation tasks.
- **Large Transformer encoder model QAT with Knowledge Distillation**
  - Perform in-depth analysis of the mechanism of KD on attention recovery of quantized large Transformer encoders.
  - Analyze quantization effect on attention behavior of transformer over various language understanding tasks.
  - Propose a new KD method and unification of multiple KD loss function to address task-dependent preference.
  - Achieve state-of-the-art language understanding accuracy for QAT with sub-2bit weight quantization for large Transformer encoder models.
- **Improving Transformer encoder QAT convergence & accuracy of few-sample fine-tuning**
  - Propose a proactive Teacher Intervention KD method for fast converging QAT of low precision pre-trained Transformers.
  - Develop gradual intervention mechanism to stabilize the recovery of subsections of Transformer layers from quantization.
  - Achieve higher accuracy of language understanding task within 12.5x shorter fine-tuning time.

SCHOLARSHIP AND AWARD

- **Integrated PhD Course Scholarship**, Full Tuition, *Hanyang University* — Spring 2021 - Spring 2024
- **Reseach Scholarship** USD 8K per year, *IoT System Semiconductor Research Center* — Spring 2021 - Spring 2023
- **AI Grand Challenge**, *Korea Ministry of Science and ICT* — Fall 2020
  - First place award in Model Compression Track
  - Compress YOLOV5s Object Detection model with 4x speed up

SKILLS

- **Programming Languages**: Python, C, C++
- **Teaching Assistant**: SOC design (Spring 2021), Introduction to SW Optimization (Fall 2023)
- **DL Frameworks**: Pytorch, Huggingface
- **Cloud Computing Platform**: NAVER NSML Machine Learning platform, KT Genie Mars Server Platform
- **English**: Served in the US Army as a representative of the Korean Army soldier for 21 Month (Jul 2017 - Apr 2019), Certified Air Traffic Control (ATC) Operator.